

CSE427s - 3 [P] MR3: Max  
Number of ...

100% (2/2)

- ✓ 1. The default behaviour in a MapReduce execution is that each Reduce task (Reducer) executes the `reduce()` function exactly once.

☐ A True  
☒ B False

- ✓ 2. *Maximum parallelism* (executing only one `reduce()` function per Reduce task) is a good idea.

☐ A Yes  
☒ B No

- ⊘ 3. Give **one argument** why maximum parallelism (executing only one `reduce()` function per Reduce task) is not good idea.

Because it will take a long time to transfer the intermediate files. If we have 1 reducer, a partition in a node will be read and transferred to one reducer. If we have  $n$  reducers that take care of one key respectively, this partition will be read into different key values and transferred  $n$  times to  $n$  reducers.