

## CSE427s - 4 Lab4

80% (8/10)

- ✓ 1. **Job** is the primary interface to describe a MapReduce job to the Hadoop framework for execution.
- ☐ A True
- ☐ B False
- ✓ 2. The name of a MapReduce job is defined by the name of the driver class.
- ☐ A True
- ☒ B False
- ✓ 3. Why should you always implement ToolRunner in your MapReduce programs?
- ☐ A It speeds up the actual job execution.
- ☒ B It allows you to run your job locally using the command line.
- ☒ C It allows parameters to be passed via the command line.
- ☐ D It is required to run your job locally via Eclipse.
- ☒ E It saves time when testing different configuration settings (such as number of reducers, etc).
- ✗ 4. What are the various ways to configure MapReduce jobs? Name two possible ways. (This answer will be graded by the TA.)
- command line configuration
- configuration within driver class
- ✓ 5. Which ways of setting job configurations overwrite others and when are default values considered? Identify the correct order according to the priority (weakest to highest).
- ☐ A command line - driver - Hadoop configuration files - defaults
- ☐ B defaults - command line - driver - Hadoop configuration files
- ☐ C driver - command line - Hadoop configuration files - defaults
- ☒ D defaults - Hadoop configuration files - command line - driver
- ✓ 6. Files provided via DistributedCache are stored in HDFS in the same way as other data (i.e., when executing `hadoop fs -put`).
- ☐ A True
- ☒ B False

- ✓ 7. Which data can be provided via **DistributedCache**?
- ☒ A look-up tables
  - ☐ B job statistics (that will be updated by the currently running job)
  - ☐ C detailed transaction or customer information for every data record
  - ☒ D java libraries that are used in Mappers or Reducers
  - ☐ E entire log file dumps
- ✓ 8. The output of each Reduce Task is re-sorted by the Reducer output **key**.
- ☐ A True
  - ☒ B False
- ✗ 9. In the output of a MR job the **values** are sorted.
- ☒ A True
  - ☐ B False
- ✓ 10. In a secondary sort implementation, which key is used to determine which Reducer an entry goes to?
- ☒ A primary key
  - ☐ B secondary key
  - ☐ C composite key
- ⊘ 11. Why do we need to implement the **sort comparator** for the secondary sort example in this lab? (This question will not be graded, so you may answer it for yourself after finishing the lab.)
- to be answered
- ⊘ 12. Why do we need to implement the **group comparator** for a secondary sort such as the one discussed in this lab? (This question will not be graded, so you may answer it for yourself after finishing the lab.)
- to be answered